

DPhil: Transfer of Status Literature Review

Shaan Desai

September 20, 2020

Contents

1	Introduction	2
2	Artificial Intelligence	2
3	Inductive Biases	3
3.1	Graph Neural Networks	3
3.2	Integrative Biases	4
3.3	Physics priors	5
3.4	Gradient Learning	5
3.5	Physics Informed Neural Networks	6
3.6	Energy Conserving Networks	7
3.6.1	Deep Lagrangian Network	7
3.6.2	Hamiltonian Neural Networks	7
3.6.3	Variational Integrator Networks	8
3.6.4	Deep Energy Network	9
3.6.5	Unsupervised Learning of Lagrangian Dynamics	9
3.6.6	Lagrangian Neural Networks	10
3.6.7	Modeling System Dynamics with PINNs on Lagrangian Mechanics	10
3.7	Symplectic Networks	11
3.7.1	Symplectic ODE Net	11
3.7.2	Symplectic Recurrent Neural Network	11
3.7.3	Deep Hamiltonian Networks based on symplectic integrators	11
3.7.4	SympNets	11
3.8	Gaussian-Based Physics Networks	12
3.8.1	Learning Constrained Dynamics with Gauss' Principle ad- hering Gaussian Processes	12
3.8.2	Bayesian Hidden Physics Models	12
3.9	Data Driven Model Discovery	13
3.9.1	Discovering Physical Concepts with Neural Networks . . .	13
3.9.2	Discovering Governing equations from data by sparse iden- tification of nonlinear dynamical systems	13

3.9.3 Symbolic Pregression	14
3.10 Causality	14

1 Introduction

The past decade has seen an immense dedication to research in Artificial Intelligence, most notably in Deep Neural Networks (DNNs). DNNs have been used across a whole host of domains including image classification, machine translation and time-series forecasting. Despite these grand successes, DNNs have numerous limitations. In particular, DNNs are black boxes meaning they aren't interpretable, they require many datapoints (> 5000 points) to learn and are highly sensitive to hyperparameter tuning. Given such limitations, the physics community has been apprehensive about its adoption. However, over the past 2 years, significant research on inductive biases has paved a new road to embed knowledge from physics into DNNs. This prior knowledge has been shown to make DNNs more data efficient and more interpretable. In this review, we look at state-of-the-art methods at the intersection of deep networks and physics. For every method, we summarize the key tricks that drive performance and also mention possible pain points worth investigating further.

2 Artificial Intelligence

Humans have long tried to understand how the mind works. From a biological perspective it is clear that connectomics paves a pathway to understanding how we think. With the advent of computers and 'large' processing power, many researchers began to ask whether the human mind can be replicated using a computer. Early work by walter and warren was the first to point out a means to build neural network logic. There has been a slew of advents in neural processing, convnets, rnns, graphs, tensor nets, deep RL. All seeking to solve more tasks/increase generalizability. To date, there has been no general purpose solution for a vast array of tasks e.g. a human can learn to create art, sit math exams, play sports, yet a robot so far can only do one of these tasks well. Numerous limitations in terms of processing power and memory access still exist. Despite these limitations, neural networks have brought about significant advancements to society. In what follows, we outline some of the major advances, those which have been fundamental pillars to research conducted during this phd. We follow no particular order.

3 Inductive Biases

An inductive bias in machine learning is prior information used to guide model building. In simple linear regression, for example, we assume the distribution of the noise in our data follows a Gaussian. This is a natural inductive bias because prior to seeing any data, we have assumed our noise model follows this form.

In neural networks, architectures with variable depth, form and activations for example, serve as inductive biases. In essence, we encode initial assumptions about the complexity of data via inductive biases.

A feedforward deep neural network naturally encodes ‘nonlinearity’ as a prior and arguably makes the least assumptions about the underlying data. Convolutional Neural Networks assume spatial representations can be captured. Recurrent Neural Networks assume temporal data as inputs.

In what follows, we highlight some of the recent advances in inductive biases, particularly as they relate to solving problems in physics.

3.1 Graph Neural Networks

The state of a physical system can be represented by a graph $G = (u, V, E)$ [3]. For example, a node (V) can be used to represent a particle in an N-body problem. These nodes can be used to represent the core features of the particle e.g. position, momentum, mass and particle constants. Edges (E) can represent forces between the particles and ‘Globals’ (u) can represent constants such as air density, the gravitational constant etc. We note that graphs can be fully connected or sparse, depending on the specific use case. In representing physical systems this way, we impart structure on our data which forms an important inductive bias when learning physics [4, 3, 27, 28, 9, 29, 26, 16, 8]. Representations of this form allow us to carry out learning in multiple complex domains because such graphs can be input to Graph Neural Networks which can be used to update the graph attributes. In addition, graphs allow for increased interpretability [3].

Some major applications include:

- Interaction-physics based problems naturally benefit from graphs [3].
- We see that graph networks trained to learn Hamiltonians achieve great results in rolling out trajectories of N-body systems [25].
- Graphs have shown significant promise in explaining the phase transition of glassy materials [2].
- They’re even shown to model complex fluid like systems from visual data [26].

Graph Neural Networks are considered to impart a relational inductive bias, as convolutional networks are considered to impart a spatial bias and recurrent networks to impart a temporal bias.

Inspired by this work we see two major steps in moving this research forward. Firstly, their use in physics opens up a tried-and-tested pathway to solve more complex problems particularly in accounting for material interactions. As such, we see good scope to use graph networks for large interacting systems as has been shown in [26]. Secondly, most systems to date have looked at graphs

for classical physics but literature from 2004 suggests graphs, inherent in their relational structure, can also capture Ising-like hamiltonian structure. In addition, although Hamiltonian Graph Networks show strong performance, not much work has been done to uncover the learned relations and whether they are consistent with ground truth interactions.

3.2 Integrative Biases

In 2018, NeuralODE brought to light a meaningful connection between residual networks and integrative steps. For example, if we use a neural network to learn a function F then use a residual block to add this function multiple times, in the continuous limit, this approximates a differential equation.

Residual networks consist of a discrete set of steps:

$$h_{t+1} = h_t + f(h_t, \theta_t)$$

which take on the form a discrete euler equation. In the continuous limit, this becomes:

$$\frac{dh(t)}{dt} = f(h(t), t, \theta)$$

We can therefore parametrize any differential equation as a neural network and integrate it as long as we know its initial condition and final time step evaluation. However, one of the challenges of passing a neural network output to an integrator is the fact that the integrator induces additional operations on the weights. Neural ODEs via the adjoint method allow us to integrate this neural network with constant memory cost [6].

The integrative bias here is used in numerous systems aimed at learning continuous time dynamics. In many settings a one-step integration is applied. For this, we do not need to use any complex ODESolver. Instead, a simple 'manual' integration is used. There are indeed settings where a multi-step integration is used, such as in [36] and in [24].

[10] identifies the instability of using the adjoint method to compute continuous depth networks and proposes a method to resolve this issue.

Extensions of this work have been adopted across numerous methods we are about to discuss. Some challenges involve identifying the right integrator of choice e.g. symplectic integrators for symplectic flows. Some additional methods have proposed a separate integrator built on neural networks to avoid the continuous depth neural networks.

3.3 Physics priors

Broadly, the intersection of physics and AI falls into one of two domains, physics for AI or AI for physics. The former uses techniques from physics to develop

and improve learning algorithms in general. The latter uses existing learning approaches (with adaptations) to predict physics, for example ML for materials. In this review, we focus our attention on the latter.

Physicists have long been interested in using learning tools to predict physics based systems. Some of these include predicting magnetic properties of 2-D materials, predicting the time evolution of N-body systems and even using AI to understand phase transitions. However, numerous challenges still remain in terms of data-efficient learning, reducing computational cost, improving predictive accuracy and learning better representations of the underlying physical process. Researchers have identified methods to address these challenges, but arguably the most promising hinges on physics-informed priors embedded in learning. It has been shown that models enriched with physically-informed priors i.e. models which consist of some knowledge about the physical system apriori, significantly outperform traditional methods in terms of data-efficiency and predictive accuracy. This has sparked a sharp interest in building both task-specific and general physics priors to improve learning. In this section, we summarize some of the core developments over time in physics informed inductive biases.

3.4 Gradient Learning

Although most modern methods cite gradient learning by Witkoskie [32] as one of the earliest efforts designed to improve learning of physics in neural networks, we actually find that more sophisticated approaches were developed prior to this effort.

In 1996 James Howse [13] presented a paper which highlights a few things on identifying dynamical systems. The paper introduces a model inspired by defining a generalized functional form for dynamical systems. Using a potential $V(x)$ we can partition an n-dimensional phase space in which the first space is normal to a level surface $V(x) = k$, and the second is tangent to $V(x) = k$. Systems which always move downhill are gradient like systems: $\dot{x} = -P(x)\nabla_x V(x)$. Systems which remain at constant potential are Hamiltonian like: $\dot{x} = Q(x)\nabla_x V(x)$. The combination of the two can be used to get total dynamics.

The entire framework is a parameter fitting one because we have a sense for the prior functional form of the dynamics.

In addition, the notion of embedding physically-informed inductive biases in neural networks can be found in numerous early work aimed at modelling materials [32, 20, 30, 23, 33]. For example, early efforts by Witkoskie and Doren [32] demonstrate that in contrast to directly learning a potential energy surface, the inclusion of gradient learning can drive a network to accurately model the forces.

Therefore, the loss function takes the form:

$$\left\| \begin{bmatrix} \hat{y} \\ \hat{\dot{y}} \end{bmatrix} - \begin{bmatrix} y \\ \dot{y} \end{bmatrix} \right\|_2^2$$

This addition means that we can supplement the learning process with additional information and hence improve the learnt potential surface. The fundamental idea being that if we have access to supplemental data such as the gradients, but fewer data points, we might actually learn a surface with higher accuracy than if we had many data points with no gradient information.

This result inspires us to look more closely at combining inductive biases. Namely, by adding well known priors together, we may make learning data-efficient and more accurate.

3.5 Physics Informed Neural Networks

PINNs is one of the first methods to use the backpropagation technique from neural networks to actually compute the gradients of a function with respect to the inputs of that given function. For example, given t and x , PINNs compute $u(t, x)$ through which they can compute gradients of u w.r.t. t and x . A concrete example is this differential equation:

$$u_t + N[u] = 0$$

where N is a nonlinear differential operator and $u(t, x)$ denotes the latent hidden solution.

Using the equation we can define:

$$f = u_t + N[u]$$

then we can, for any set (t, x) , compute $u(t, x)$ with a neural network. Using backpropagation, we can then compute partial derivatives w.r.t the input variables. The final results can thus be stored in f . A simple L2 loss to minimize the predicted u vs the ground truth u (for initial and boundary conditions) coupled with a penalization of the function f (since it should be zero) at specific collocation points results in PINNs.

PINNs are used for both data-drive solution and data-driven discovery. In other words, PINNs can be used to identify a systems governing equations or given a set of equations, be used to continuously solve the system given an initial condition and points of evaluation.

One limitation to this approach is having to know the functional form of the equation. However, the authors of this paper also developed deep hidden physics

models [21] that are designed to learn the functional form without any prior knowledge of the function.

3.6 Energy Conserving Networks

3.6.1 Deep Lagrangian Network

DeLAN net is the first network to use a Lagrangian embedded in a neural network to learn the dynamics of a system.

In the paper they define:

$$L = T - V$$

and

$$\frac{d}{dt} \frac{\partial L}{\partial \dot{q}} - \frac{\partial L}{\partial q} = \tau$$

where τ represents generalized forces.

If we are dealing with a rigid body then: $T = \dot{q}^T M(q) \dot{q}$ where M is the inertia matrix. Replacing this equation in euler-lagrange results in:

$$\frac{d}{dt}(M(q)\dot{q}) - \frac{\partial V}{\partial q} = 0$$

$$M(q)\ddot{q} = \dot{M}(q)\dot{q} + \frac{\partial V}{\partial q}$$

To get \ddot{q} all we need is M and its time derivative.

$$p = M(q)\dot{q}$$

As a result of this framing, DeLAN net uses multiple network heads to compute individual components of the equation before combining them.

3.6.2 Hamiltonian Neural Networks

ing neural networks to accurately learn classical dynamics from data problems has been Hamiltonian Neural Networks [12]. It has been demonstrated that dynamic predictions through time can be improved using Hamiltonian Neural Networks (HNNs) which endow models with a Hamiltonian constraint. The Hamiltonian is an important representation of a dynamical system because it is one of two approaches that generalize classical mechanics. The Hamiltonian \mathcal{H} is a scalar function of position $\mathbf{q} = (q_1, q_2, \dots, q_M)$ and momentum $\mathbf{p} = (p_1, p_2, \dots, p_M)$. In representing physical systems with a Hamiltonian, one can simply extract the

time derivatives of the inputs by differentiating the Hamiltonian with respect to its inputs (see Eqn. 1.)

$$\frac{d\mathbf{q}}{dt} = \frac{\partial \mathcal{H}}{\partial \mathbf{p}}, \quad \frac{d\mathbf{p}}{dt} = -\frac{\partial \mathcal{H}}{\partial \mathbf{q}} \quad (1)$$

As a consequence, it is noted in [12] that by accurately learning a Hamiltonian, the system’s dynamics can be naturally extracted through backpropagation. This information allows us to build two 1st-order differential equations which can be used to update the state space, (\mathbf{q}, \mathbf{p}) . Equation 2 shows this integral, in which we define the symplectic gradient $\mathbf{S} = \left[\frac{\partial \mathcal{H}}{\partial \mathbf{p}}, -\frac{\partial \mathcal{H}}{\partial \mathbf{q}} \right]$:

$$(\mathbf{q}, \mathbf{p})_{t+1} = (\mathbf{q}, \mathbf{p})_t + \int_t^{t+1} \mathbf{S}(\mathbf{q}, \mathbf{p}) dt \quad (2)$$

It can be shown that the Hamiltonian in many systems also represents the total energy of the system. Therefore, the Hamiltonian is a powerful inductive bias that can be utilised to evolve a physical state while maintaining energy conservation.

As we will discuss later, HNNs are great at learning in low-dimensional settings for a few integration steps. However, scaling this network to more challenging problems proves difficult. Fortunately, many alternatives have been proposed.

One limitation we hope to investigate further in this domain is embedding the energy term. HNNs appear to learn a scaled/transformed energy value. From preliminary experiments we find that penalizing the loss function with an energy term actually hurts the learning process indicating instability issues with the approach.

3.6.3 Variational Integrator Networks

Lagrangian mechanics offers an alternative to the Hamiltonian in generalizing a dynamical system. Rather than position and momentum (canonical coordinates) defining the state space, Lagrangian mechanics is defined using a generalized coordinate state space $(\mathbf{q}, \dot{\mathbf{q}})$. This is particularly useful in physical settings where the description and measurement of generalized coordinates may be easier to work with than canonical coordinates [18]. Given these coordinates, Joseph-Louis Lagrange showed that a scalar value \mathcal{A} , referred to as the action, can be defined as the integral of a Lagrangian, $\mathcal{L}(\mathbf{q}, \dot{\mathbf{q}})$:

$$\mathcal{A} = \int_t^{t+1} \mathcal{L}(\mathbf{q}, \dot{\mathbf{q}}) dt \quad (3)$$

The integral can be thought as inducing multiple paths between points in state space i.e. multiple walks in the domain of $(\mathbf{q}, \dot{\mathbf{q}})$. However, only one path is a stationary state of the action integral. This state lets us move from $t \rightarrow t + 1$

with minimal energy. It can be shown, through variational calculus, that this stationary state must satisfy the Euler-Lagrange equation:

$$\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) = \frac{\partial \mathcal{L}}{\partial \mathbf{q}} \quad (4)$$

Although complex in form, the action integral and the Euler-Lagrange equations can be discretized and collectively form the basis for variational integrators. The work in [24] shows that, by adopting this approach, one can develop VINs which make network learning in noisy data-settings more robust. Similar to Hamiltonians, Lagrangians in classical mechanics are also connected to the kinetic energy \mathcal{T} and potential energy \mathcal{V} via:

$$\mathcal{L} = \mathcal{T}(\mathbf{q}, \dot{\mathbf{q}}) - \mathcal{V}(\mathbf{q}, \dot{\mathbf{q}}) \quad (5)$$

Furthermore, Variational Integrators are symplectic and momentum conserving [17].

The paper introduces the importance of symplectic integrators but does not discuss how to improve accuracy and scale up to higher dimensions. It is with this in mind that VIGNs was borne.

3.6.4 Deep Energy Network

A new method is introduced to tackle the challenge of learning energy conserving physics more broadly. The method claims to be more general as it can learn from PDEs as well as ODEs, and it avoids discretisation errors that arise from using RK integrators. The method does this by a) introducing a general formalism which highlights how the rate of change of states can be cast into a matrix G times the gradient of a Hamiltonian and b) by discretising the gradient of the hamiltonian with a frechet derivative.

$$\dot{u} = G \nabla H \quad (6)$$

The main benefit is that it can solve a broader class of problems compared to current solutions in the field by modelling an additional G matrix e.g. Friction systems (ODE and PDE) Discrete PDEs (PDE) Maintains energy and mass conservation via discretisation which preserves the geometric structure (a.k.a volume preservation) Good form for discrete auto differentiation Great PDE dataset motivations Can extend research to identify discrete PDEs

3.6.5 Unsupervised Learning of Lagrangian Dynamics

In [37], the authors show that the full Lagrangian dynamics can be learned from visual data. The paper introduces a co-ordinate aware VAE to encode the latent space and then computes the derivatives of the state. The introduction of co-ordinate awareness is to ensure a bijective function. The paper shows that the motion of a pendulum can be learnt from visual data. More importantly, via energy shaping, the pendulum can be constrained into a configuration q^* of choice.

3.6.6 Lagrangian Neural Networks

The main premise of LNNs [8] is to tackle the problem of dealing with coordinate spaces. Many datasets do not usually consist of canonical position and momentum, rather they use generalized coordinates. As such, LNNs aim to tackle learning from generalized coordinates. In addition, they provide a more general framework than DeLaNs which were designed to work well with continuous control applications. The unique trick that LNNs introduce over other methods is that they do not assume any form for the lagrangian.

In vectorized form, euler lagrange is:

$$\frac{d}{dt}\nabla_{\dot{q}}\mathcal{L} = \nabla_q\mathcal{L} \quad (7)$$

Using chain rule to expand the time derivative e.g. allow $\nabla_{\dot{q}}\mathcal{L}$ to be a function of q and \dot{q} then:

$$(\nabla_{\dot{q}}\nabla_{\dot{q}}^T\mathcal{L})\ddot{q} + (\nabla_q\nabla_{\dot{q}}^T\mathcal{L})\dot{q} = \nabla_q\mathcal{L} \quad (8)$$

Then, with matrix inversion, one can obtain \ddot{q} .

Using this term, one can minimize the loss on the state vector $[q, \dot{q}, \ddot{q}]$. The paper shows promising results on the double pendulum, relativistic particle in a uniform potential and on the wave equation. Upon inspection of the training scheme for the highly sensitive double pendulum, we do find that training times and data points are quite large, creating a space to find a more optimal learning scheme.

3.6.7 Modeling System Dynamics with PINNs on Lagrangian Mechanics

Unlike LNNs, this paper introduces a functional form for the Lagrangian. The Lagrangian in this paper [22] is assumed to be $L = T - V$. In addition, they introduce non conservative forces and the final form is:

$$M(q)\ddot{q} + C(q, \dot{q})\dot{q} + G(q) = Q^{ncons}$$

The approach taken by this paper is to feed in the respective components e.g. q, \dot{q} into separate neural network heads designed to predict one of M, C and G. Using this technique, \ddot{q} is backed out and then fed to an RK-4 integrator.

The approach is tethered to a more traditional way of using NNs for predictions and is one of the reasons why model performance is not as good as it can be if backpropagation was used to compute the second derivative of q . In addition, the wrong choice of integrator is a bottleneck for long range predictions in this setting.

3.7 Symplectic Networks

3.7.1 Symplectic ODE Net

Symp-ODEN [36] extends hamiltonian neural networks into the control domain. If external control is affine and influences the change in generalized momenta then:

$$\begin{bmatrix} \dot{q} \\ \dot{p} \end{bmatrix} = \begin{bmatrix} \frac{\partial H}{\partial p} \\ -\frac{\partial H}{\partial q} \end{bmatrix} + \begin{bmatrix} 0 \\ g(q) \end{bmatrix} u$$

If $\text{rank}(g(q))$ is $\text{rank}(q)$ the system is fully actuated. For such systems, a controller $u = \beta(q) + v(p)$ can be designed to change the potential energy landscape so as to force a system toward a specific configuration.

If the hamiltonian is a newtonian then -

$$\beta(q) = g^T (g g^T)^{-1} (\partial V / \partial q - \partial V_d / \partial q)$$

Arguably the first paper to think extensively about the bottlenecks presented in HNNs. The paper brings to light the need for symplecticity, control inputs and a proper use of NeuralODE.

3.7.2 Symplectic Recurrent Neural Network

Symplectic RNN [7] makes two major contributions to the physics neural network space. Firstly, they take Hamiltonian Neural Networks and convert the integration scheme from euler to a leapfrog. Secondly, they adopt a NeuralODE style integrator to integrate across multiple timesteps from an initial condition. However, the results in this paper overlap greatly with Symp-ODEN.

3.7.3 Deep Hamiltonian Networks based on symplectic integrators

This paper [38] reviews the integrator of choice for hamiltonian neural networks. One of its primary objectives is to establish the difference in performance between using a symplectic vs non-symplectic integrator from a theoretical perspective. Quite evidently, using a symplectic integrator allows us to conserve the phase-space volume of the system - a crucial component in preserving energy. Their results reiterate the need for symplectic integrators when dealing with hamiltonian-like systems.

3.7.4 SympNets

A brand new class of methods is proposed . It avoids the need for a separable Hamiltonian and more importantly, is designed to eliminate the need for back-propagating the Hamiltonian with respect to the input. In this framework, a sequence of symplectic maps are used to transform the input into the output.

The symplectic map can be split into an upper triangular matrix and lower triangular matrix with diagonals set to 1. Non-diagonal terms are parametrized by a NN and by stacking a range of symplectic maps together, one can learn dynamics better. In fact, the results show that by doing this, the learned trajectory is significantly more accurate than HNNs with symplectic integrators!

3.8 Gaussian-Based Physics Networks

3.8.1 Learning Constrained Dynamics with Gauss' Principle adhering Gaussian Processes

The paper [11] introduces GP model that utilises mechanical constraints as prior knowledge for learning dynamics of systems. The GP is constrained to satisfy Gauss Principle.

Using UKE, the acceleration of a particle can be disentangled into unconstrained acceleration, ideal part of constraint and non-ideal part of constraint. Using this knowledge, one can feed this structural form into the mean of a GPR. The process can be used to set priors on unknown parameters. The learned model appears to capture dynamics much better than standard GPs without the underlying equation needed to be satisfied in the mean of the GP.

3.8.2 Bayesian Hidden Physics Models

In [1], the author presents a novel approach to tackle learning from noisy data. The paper is arguably the first method to combine gaussian processes with physics informed neural networks in a unified manner.

Simply put, the learning takes place in 3 stages. The first is, given inputs x and t , compute $u(x, t)$ using a neural network. The log likelihood is defined as:

$$L_i^u = \log p(D_i^u | \theta_i^u) = \sum_{j=1}^{n_{st}} \log p(\hat{u}_j^i | u(x_j, t_j; \theta_i))$$

Then, using the computed u , define a prior over the derivatives of u . The purpose of doing this is for data-driven discovery. In other words, rather than computing k partial derivatives of u w.r.t the input x , we place a gaussian process prior on the k partial derivatives.

$$L^f = p(\hat{u}_t | V) = \mathcal{N}(\hat{u}_t | \mu(V), K_f + \sigma_f^2 I)$$

As such, one can obtain the posterior using variational inference since we do not know the distribution of L^f .

The main breakthrough of using this approach is uncertainty quantification of the learned operators as well as the convergence in v-space the method illustrates.

3.9 Data Driven Model Discovery

3.9.1 Discovering Physical Concepts with Neural Networks

Although this review has looked closely at embedding physical laws into neural networks, some approaches such as [14] take on a different approach of not assuming any prior. The paper sets the 'scientific process' as the prior. By feeding inputs into an encoder to learn a representation and then querying this representation before decoding, the experimental process of learning laws is designed to become a part of the network.

The only caveat to this modelling is the use of disentangled VAEs that are designed to produce orthogonal representations. The assumption here being that every new variable added to the representation should be independent and alter the output landscape separately i.e. no linear combinations/the variables should span the output space.

Results illustrate that such a system can learn accurate latent representations necessary to evolve the time dynamics of a system, including conservation laws. The network also learns how to switch co-ordinate systems and the dimension of underlying quantum systems.

This presents an alternative view to learning physical systems. Motivated by this result, one can easily draw a comparison across methods to illustrate how much more data/time is needed to establish the right representation in comparison to methods naturally designed to embed physical laws.

3.9.2 Discovering Governing equations from data by sparse identification of nonlinear dynamical systems

The proposed method in [5] tackles learning the governing equations of a physical system by inspecting the sparsity of a function f that satisfies:

$$dx/dt = f(x(t))$$

For most systems, the paper identifies that only a small group of derivatives are relevant to compute f .

The method first computes higher order polynomials of the inputs and then uses sparse regression to select the parameters.

This approach of 'selecting' from a bag of potential functions has been explored extensively in determining governing equations. However, the method relies on having to explicitly build a large matrix of permutations.

3.9.3 Symbolic Pregression

Fundamental early work aimed at devising unique ways to find regression coefficients using a supervised learning framework. Symbolic pregression combines this approach with that of the hamiltonian neural net to build an unsupervised framework for discovery of differential equations [31]. The paper still does not reveal the functional form of the underlying constraint but the results look deeply promising in dealing with visual data.

3.10 Causality

Bayesian networks are probabilistic graphical models that can represent conditional dependencies using a Directed Acyclic Graph. The main logic behind these methods is that a probabilistic approach to determining a marginal distribution can be linked to a graph structure if certain conditional dependencies are met.

Learning DAGs from data is an NP-hard problem, but the NOTEARS algorithm [35] presents a way to tackle this combinatorial search by converting the problem into a continuous optimization problem. The main contribution of the paper is a penalization of the weight matrix so that it is acyclic. From studies of graphs we know that the n 'th power of an adjacency matrix gives us the lengths of walk k between two nodes. As such, all we need to do is enforce that all the diagonals of the powers of the adjacency matrices are set to zero.

In other words:

$$\text{tr}(I - B)^{-1} = d$$

must be satisfied. Embedding this into the learning process is imparting an inductive bias on learning.

It was with this in mind that DYNOTEARS [19] were able to extend the work into the temporal setting and learn dynamic bayesian networks. In principle, by adding an additional weight matrix to learning and penalizing it in the same way that NOTEARS does.

Furthermore, it has been shown that the NOTEARS constraint can be embedded in graph neural networks [15, 34].

References

- [1] S. Atkinson. Bayesian Hidden Physics Models: Uncertainty Quantification for Discovery of Nonlinear Partial Differential Operators from Data. *arXiv:2006.04228 [cs, stat]*, June 2020. arXiv: 2006.04228.
- [2] V. Bapst, T. Keck, A. Grabska-Barwińska, C. Donner, E. D. Cubuk, S. S. Schoenholz, A. Obika, A. W. R. Nelson, T. Back, D. Hassabis, and P. Kohli.

- Unveiling the predictive power of static structure in glassy systems. *Nature Physics*, 16(4):448–454, Apr. 2020. Number: 4 Publisher: Nature Publishing Group.
- [3] P. W. Battaglia, J. B. Hamrick, V. Bapst, A. Sanchez-Gonzalez, V. Zambaldi, M. Malinowski, A. Tacchetti, D. Raposo, A. Santoro, R. Faulkner, C. Gulcehre, F. Song, A. Ballard, J. Gilmer, G. Dahl, A. Vaswani, K. Allen, C. Nash, V. Langston, C. Dyer, N. Heess, D. Wierstra, P. Kohli, M. Botvinick, O. Vinyals, Y. Li, and R. Pascanu. Relational inductive biases, deep learning, and graph networks. *arXiv:1806.01261 [cs, stat]*, Oct. 2018. arXiv: 1806.01261.
 - [4] P. W. Battaglia, R. Pascanu, M. Lai, D. Rezende, and K. Kavukcuoglu. Interaction Networks for Learning about Objects, Relations and Physics. *arXiv:1612.00222 [cs]*, Dec. 2016. arXiv: 1612.00222.
 - [5] S. L. Brunton, J. L. Proctor, and J. N. Kutz. Discovering governing equations from data by sparse identification of nonlinear dynamical systems. *Proceedings of the National Academy of Sciences*, 113(15):3932–3937, Apr. 2016.
 - [6] R. T. Q. Chen, Y. Rubanova, J. Bettencourt, and D. K. Duvenaud. Neural Ordinary Differential Equations. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems 31*, pages 6571–6583. Curran Associates, Inc., 2018.
 - [7] Z. Chen, J. Zhang, M. Arjovsky, and L. Bottou. Symplectic Recurrent Neural Networks. *arXiv:1909.13334 [cs, stat]*, Apr. 2020. arXiv: 1909.13334.
 - [8] M. Cranmer, S. Greydanus, S. Hoyer, P. Battaglia, D. Spergel, and S. Ho. Lagrangian Neural Networks. *arXiv:2003.04630 [physics, stat]*, Mar. 2020. arXiv: 2003.04630.
 - [9] M. D. Cranmer, R. Xu, P. Battaglia, and S. Ho. Learning Symbolic Physics with Graph Networks. *arXiv:1909.05862 [astro-ph, physics:physics, stat]*, Nov. 2019. arXiv: 1909.05862.
 - [10] E. Dupont, A. Doucet, and Y. W. Teh. Augmented Neural ODEs. *arXiv:1904.01681 [cs, stat]*, Oct. 2019. arXiv: 1904.01681.
 - [11] A. R. Geist and S. Trimpe. Learning Constrained Dynamics with Gauss Principle adhering Gaussian Processes. *arXiv:2004.11238 [cs, eess, stat]*, Apr. 2020. arXiv: 2004.11238.
 - [12] S. Greydanus, M. Dzamba, and J. Yosinski. Hamiltonian Neural Networks. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d. Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems 32*, pages 15379–15389. Curran Associates, Inc., 2019.

- [13] J. W. Howse, C. T. Abdallah, and G. L. Heileman. Gradient and Hamiltonian Dynamics Applied to Learning in Neural Networks. In D. S. Touretzky, M. C. Mozer, and M. E. Hasselmo, editors, *Advances in Neural Information Processing Systems 8*, pages 274–280. MIT Press, 1996.
- [14] R. Iten, T. Metger, H. Wilming, L. del Rio, and R. Renner. Discovering physical concepts with neural networks. *Physical Review Letters*, 124(1):010508, Jan. 2020. arXiv: 1807.10300.
- [15] S. Lachapelle, P. Brouillard, T. Deleu, and S. Lacoste-Julien. GRADIENT-BASED NEURAL DAG LEARNING. page 23, 2020.
- [16] L. Lamb, A. Garcez, M. Gori, M. Prates, P. Avelar, and M. Vardi. Graph Neural Networks Meet Neural-Symbolic Computing: A Survey and Perspective. *arXiv:2003.00330 [cs]*, Mar. 2020. arXiv: 2003.00330.
- [17] A. Lew, J. E. Marsden, M. Ortiz, and M. West. AN OVERVIEW OF VARIATIONAL INTEGRATORS. *Finite Element Methods*, page 18.
- [18] J. E. Marsden and M. West. Discrete mechanics and variational integrators. *Acta Numerica*, 10:357–514, May 2001.
- [19] R. Pamfil, N. Sriwattanaworachai, S. Desai, P. Pilgerstorfer, P. Beaumont, K. Georgatzis, and B. Aragam. DYNOTEARS: Structure Learning from Time-Series Data. *arXiv:2002.00498 [cs, stat]*, Apr. 2020. arXiv: 2002.00498.
- [20] A. Pukrittayakamee, M. Malshe, M. Hagan, L. M. Raff, R. Narulkar, S. Bukkapatnum, and R. Komanduri. Simultaneous fitting of a potential-energy surface and its corresponding force fields using feedforward neural networks. *The Journal of Chemical Physics*, 130(13):134101, Apr. 2009.
- [21] M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, Feb. 2019.
- [22] M. A. Roehrl, T. A. Runkler, V. Brandtstetter, M. Tokic, and S. Obermayer. Modeling System Dynamics with Physics-Informed Neural Networks Based on Lagrangian Mechanics. *arXiv:2005.14617 [cs, stat]*, May 2020. arXiv: 2005.14617.
- [23] M. Rupp, A. Tkatchenko, K.-R. Müller, and O. A. von Lilienfeld. Fast and Accurate Modeling of Molecular Atomization Energies with Machine Learning. *Physical Review Letters*, 108(5):058301, Jan. 2012. Publisher: American Physical Society.
- [24] S. Saemundsson, A. Terenin, K. Hofmann, and M. P. Deisenroth. Variational Integrator Networks for Physically Meaningful Embeddings. *arXiv:1910.09349 [cs, stat]*, Oct. 2019. arXiv: 1910.09349.

- [25] A. Sanchez-Gonzalez, V. Bapst, K. Cranmer, and P. Battaglia. Hamiltonian Graph Networks with ODE Integrators. *arXiv:1909.12790 [physics]*, Sept. 2019. arXiv: 1909.12790.
- [26] A. Sanchez-Gonzalez, J. Godwin, T. Pfaff, R. Ying, J. Leskovec, and P. W. Battaglia. Learning to Simulate Complex Physics with Graph Networks. *arXiv:2002.09405 [physics, stat]*, Feb. 2020. arXiv: 2002.09405.
- [27] A. Sanchez-Gonzalez, N. Heess, J. T. Springenberg, J. Merel, M. Riedmiller, R. Hadsell, and P. Battaglia. Graph networks as learnable physics engines for inference and control. *arXiv:1806.01242 [cs, stat]*, June 2018. arXiv: 1806.01242.
- [28] S. Seo and Y. Liu. Differentiable Physics-informed Graph Networks. *arXiv:1902.02950 [cs, stat]*, Feb. 2019. arXiv: 1902.02950.
- [29] S. Seo, C. Meng, and Y. Liu. PHYSICS-AWARE DIFFERENCE GRAPH NETWORKS FOR SPARSELY-OBSERVED DYNAMICS. page 15, 2020.
- [30] J. S. Smith, O. Isayev, and A. E. Roitberg. ANI-1: an extensible neural network potential with DFT accuracy at force field computational cost. *Chemical Science*, 8(4):3192–3203, 2017.
- [31] S.-M. Udrescu and M. Tegmark. Symbolic Pregression: Discovering Physical Laws from Raw Distorted Video. *arXiv:2005.11212 [physics, stat]*, May 2020. arXiv: 2005.11212.
- [32] J. B. Witkoskie and D. J. Doren. Neural Network Models of Potential Energy Surfaces: Prototypical Examples. *Journal of Chemical Theory and Computation*, 1(1):14–23, Jan. 2005.
- [33] K. Yao, J. E. Herr, D. Toth, R. Mckintyre, and J. Parkhill. The TensorMol-0.1 model chemistry: a neural network augmented with long-range physics. *Chemical Science*, 9(8):2261–2269, 2018.
- [34] Y. Yu, J. Chen, T. Gao, and M. Yu. DAG-GNN: DAG Structure Learning with Graph Neural Networks. page 10.
- [35] X. Zheng, B. Aragam, P. Ravikumar, and E. P. Xing. DAGs with NO TEARS: Continuous Optimization for Structure Learning. *arXiv:1803.01422 [cs, stat]*, Nov. 2018. arXiv: 1803.01422.
- [36] Y. D. Zhong, B. Dey, and A. Chakraborty. Symplectic ODE-Net: Learning Hamiltonian Dynamics with Control. *arXiv:1909.12077 [physics, stat]*, Sept. 2019. arXiv: 1909.12077.
- [37] Y. D. Zhong and N. E. Leonard. Unsupervised Learning of Lagrangian Dynamics from Images for Prediction and Control. *arXiv:2007.01926 [cs, eess, stat]*, July 2020. arXiv: 2007.01926.
- [38] A. Zhu, P. Jin, and Y. Tang. Deep Hamiltonian networks based on symplectic integrators. *arXiv:2004.13830 [cs, math]*, Apr. 2020. arXiv: 2004.13830.